

UNICORN VYSOKÁ ŠKOLA S.R.O.

MASTER'S THESIS

2024

Bc. Filip DITRICH

UNICORN VYSOKÁ ŠKOLA s.r.o.

Softwarový vývoj



MASTER'S THESIS

Cashless festival data analysis and analytical dashboard development

Author: Bc. Filip Ditrich

Supervisor: Mgr. Václav Alt

Prague 2024

Statutory Declaration

I hereby declare that I have written my Master's Thesis on the topic of *Cashless festival data analysis and analytical dashboard development* by myself, under the guidance of my thesis supervisor, using only the technical publications and other information sources which are all quoted in the thesis and listed in the bibliography.

I declare that artificial intelligence tools have been used only for support activities and in accordance with the principle of academic ethics.

As the author of this Master's Thesis, I also declare that in association with its writing I have not violated the copyright of any third party or parties and I am fully aware of the consequences of provisions of s. 11 et seq. of Act No. 121/2000 Coll., the Copyright Act.

Furthermore, I hereby declare that the submitted hard copy of this Master's Thesis is identical to the electronic version I have submitted.

In on

Bc. Filip Ditrich

Acknowledgements

I would like to express my sincere gratitude to my supervisor, Mgr. Václav Alt, for his guidance, support, and valuable advice throughout the process of writing this thesis. I would also like to thank my family and friends for their encouragement and understanding.



**Cashless festival data analysis
and analytical dashboard
development**

Analýza dat bezhotovostního festivalu
a vývoj analytického dashboardu



Abstract

TODO

Keywords: TODO

Abstrakt

TODO

Klíčová slova: TODO

Contents

Introduction	10
Background and Motivation	10
Problem Statement	12
Objectives of the Work	14
Scope of the Study	19
1 Data and Methodology	21
1.1 Environment and local setup	21
1.1.1 Data Obtaining and Preparation	22
1.1.2 Local Database Setup	22
1.1.3 Local Database Modifications	23
1.2 Data Anonymization	24
1.3 Data Structure	26
1.3.1 Financial Data	26
1.3.2 Customer Data	27
1.3.3 Event Data	27
1.3.4 Data Processing Views	28
1.4 Tools and Technologies	31
1.5 Conclusion	31
2 Data Analysis and Results	34
2.1 Cashflow and Revenue Sources Analysis	34
2.1.1 Chip Top-Up Analysis	35
2.1.2 Sales Analysis	37
2.1.3 Remaining Chip Balances	38
2.1.4 Total Revenue of the Organizer	40
2.1.5 Summary	43
2.2 Performance Indicators Analysis	44
2.2.1 Transactions Processing Analysis	45

2.2.2	Best Sale & Top-Up Points, Vendors, and Products Analysis	48
2.2.3	Summary	54
2.3	Beverage Consumption Analysis	55
2.3.1	Returnable Cups Analysis	55
2.3.2	Total Consumption Analysis	57
2.3.3	Popular Brands Analysis	58
2.3.4	Summary	62
2.4	Customer Analysis	62
2.4.1	Event Attendance and Timeline Analysis	63
2.4.2	Customer Segmentation	70
2.4.3	Payment Behavior Analysis	70
2.4.4	Purchase Pattern Analysis	71
3	Dashboard Implementation	72
4	Conclusion	73
4.1	Summary of Work	73
4.2	Reflections	73
4.3	Future Directions	73
	List of Figures	74
	List of Tables	75
	List of Acronyms	77
	List of Appendices	78
	Appendix A Source code of the application	78

Introduction

Background and Motivation

Payments at festivals are a crucial part of the successful event management. The shift from cash to cashless payments has been a significant trend in the last decade that has brought many benefits to both festival organizers and attendees.

However, traditional cashless payment systems utilizing payment terminals are not only expensive to reliably implement at a venue, where the internet connection is often unreliable and requires expensive Base Transceiver Stations (BTS) set-ups or in-place wired/optical internet implemented which can cost up to several million CZK, but also do not provide any insights into the data generated by the transactions. In the best case scenario, the organizers are able to generate a report of processed transactions made by each terminal.

Which, frankly, is not enough to make any actionable decisions based on the data. Moreover, given the event organizers are not data scientists, they often lack the knowledge and tools to analyze the data and extract valuable insights from it.

That is where NFCtron comes in and offers a solution that not only provides a reliable cashless payment system, with credit-based NFC chip bracelets supporting offline mode or card terminal payment solutions, but also provides a comprehensive B2B platform that allows the organizers, the vendors and event the third-party partners to benefit from the data that the system operates with.

The system is a full-scope solution that provides from initial online ticket sales and online credit top-up, through attendee check-in, on-site credit top-up attendee access control, security monitoring, vendor sales and inventory management, most importantly the fast and reliable payment processing with the real-time data analytics and reporting, all the way to the post-event automatic settlement and

reporting.

It simply provides everything an event organizer needs to successfully and efficiently manage their event without the need to worry about any technicalities or staff management. Since NFCtron does not only provide the system as a service, but also provides experienced Event Managers, cashiers, check-in brigadiers and other staff to operate the system and the event itself ensuring that the organizer can focus on the event in terms of communication, marketing, line-ups and other important aspects of the event.

Put, NFCtron offers organizers a peace of mind and a guarantee that their event will be a success.

NFCtron Company

NFCtron is a Czech company that has been operating since around 2019. In its early beginning during a COVID-19 pandemic was on the verge of survival because of the event industry being paralyzed by the government restrictions. However, the company survived and even in these difficult times managed to turn the disadvantage into an advantage by focusing on the core system and product development.

Which years later resulted in a robust and reliable system. It is now used by many event organizers across the Czech Republic and Slovakia and is currently expanding to other countries in Central Europe such as Austria, Poland and Germany. In its primary market – the Czech Republic – NFCtron penetrated the market and is now the leading cashless payment system provider for festivals and other events.

In recent years, the company has also been focusing on expanding to the Payments market focusing both on Card Acquiring and Card Issuing. From the acquiring part, it is now actively developing its own SoftPOS solution that will allow vendors to accept payments via mobile phones. On the other hand, on the Issuing part, it is also working on Card Issuing; that will allow the company to issue its own NFCtron branded payment cards in cooperation with Mastercard.

A big part of the successful market penetration was the company's focus on the business-to-business (B2B) side of the business with event organizers. Giving the organizers amounts of data improving their decision-making and providing them

with insights allowing them to optimize their events. And most importantly, providing them with economic aspects and cashflow optimizations that allows many events and festivals to survive and continue to operate.

Personal Position and Motivation

I have been with the company from the COVID-19 times. My current position is a Chief Product Officer (CPO), and I am responsible for the product development and the product management of all the products and services that NFCtron offers. This allows me to have a deep insight into the system and most importantly, access to all the data that the system generates. As previously mentioned, the company's success is based on the B2B side of the business and that means services and products provided to the event organizers.

The main product, that organizers have access to is the platform called **NFCtron Hub**. My personal goal and personal motivation to work on this thesis is to discover new ways to improve the platform and provide even more valuable insights to the organizers.

Problem Statement

Even though the system provides a lot of data, it still has a lot of potential to provide even more valuable insights to the organizers. Currently, the previously mentioned B2B platform, **NFCtron Hub**, provides real time data analytics dashboard presenting the most important KPIs and metrics to the organizers.

These KPIs and metrics summarize:

- **Total sales:** the total amount spent on online ticket sales and on-site payments.
- **Total sales in time:** the total amount above split into time intervals.
- **Total refunds:** the total number of sale reversals or refunds made including refunds from online tickets, refunds of on-site payments and chip credit refunds.

- **Chip balances:** the current balance topped-up on the NFC chip bracelets on-site or pre-topped-up online.
- **Customer orders rating:** customers can rate their orders via a NFC-tron mobile application which provides the organizers with feedback on the vendor's performance and the quality of the products sold.

Moreover, it provides less clear data such as

- **List of vendors:** the list of vendors presents at the event with their sales and rating.
- **List of products:** the list of products sold at the event with their sales and rating.
- **List of points of sale:** the list of selling points at the event with their sales and rating.
- **List of top-up places:** the list of top-up places at the event with the amount of top-ups made.
- **List of customer chips:** the list of unique individual NFC chips issued to the customers with their balance, spending and security status.
- **List of customer ratings:** the list of customer ratings with their feedback from points of sale.

And finally, it also provides unstructured data in the form of data exports in tabular format that can be used for further analysis:

- **Product exports** – list of all products sold with summarized metrics.
- **Points of sale exports** – list of all selling points with summarized metrics.
- **Vendor exports** – list of all vendors with summarized metrics.
- **Deal exports** – list of summarized sales made under a deal¹ Between the organizer and the vendor.
- **Transaction exports** – a heavy export of all transactions made at the event.

¹A deal is an arrangement between the organizer and the vendor that states the terms of the vendor's presence at the event, the products they are allowed to sell, the price of the products, the commission the vendor pays to the organizer and other terms of the deal.

- **Ticket redeems exports** – list of all tickets redeemed at the event.
- **And other exports regarding the online sales** – list of all online ticket sales, top-ups, receipts, customers and other data.

With the above giving some initial picture about the capabilities of the system, the platform (and thus the organizers) still face several problems or challenges that need to be addressed:

1. **Problem 1:** The main KPI metrics may provide some core insights, but the other less or unstructured data is not used to its potential.
2. **Problem 2:** The organizers are not data scientists and need a simple and clear way to understand the data.
3. **Problem 3:** Even with all this data, there is still a lot that can be done to dig deeper and provide more valuable insights.

Objectives of the Work

With the problems stated above, the main objective of this thesis is to analyze, answer and present results to important questions about the available data and the potential insights that can be extracted from it.

But to achieve this, it was a prerequisite to find a willing event organizer, that would provide the data and would be willing to cooperate on the project. Cooperate in terms of providing valuable insights into what they would like to know more about their event.

In Cooperation with the Event Organizer

For this purpose, I chose an undisclosed event organizer who has been a close and helpful partner of NFCtron for many seasons. Together with the organizer in the first step, we have stated the following requirements to perform the analysis:

- **Requirement 1:** The event and organizer should be kept undisclosed.

- **Requirement 2:** The data should be anonymized to not leak any possible sensitive information about vendors or customers.

The next step was to choose an event from which the data will be used. As it cannot be disclosed any further, we will refer to the event as **The Event** and the organizer as **The Organizer**.

Now the important information about **The Event** for this study is the following:

- **The Event** is a music festival that has been organized for several years now.
- **The Event** takes place in the Czech Republic at the beginning of July 2025.
- **The Event** is a 3-day event with multiple stages and multiple vendors.
- **The Event** uses NFCtron system for cashless payments and access control.
- **The Event** had around 7,000 attendees in 2024 and had a roughly 43% increase in 2025 to around 10,000 attendees.

Data Analysis Objectives

The final step was to define questions or data analysis objectives that should be answered or achieved by the end of the thesis.

Together in the cooperation with **The Organizer** and several internal colleagues in NFCtron, we have defined the following questions for the data analysis:

Cashflow and Revenue

What was the organizer's total revenue and what did it consist of?

- **RQ1:** *What was the total revenue?*
- **RQ2:** *How much and how was the balance topped up on the chips?*
- **RQ3:** *What is the remaining balance on all chips after the event and refunds?*

- **RQ4:** *What was the total sales of the event, how much of it was sold by the organizer, and how many external vendors?*

These questions should shed light on the event’s cash flow and revenue sources, which the platform does not currently cover in detail. Possible answers to these questions may provide the organizer with valuable insights into the event’s economic aspects, allowing them to optimize cashflow and revenue sources for the following event.

Performance

- **RQ5:** *How many transactions were processed in total, and when was the system’s largest “peak” in transaction volume?*
- **RQ6:** *What was the average transaction processing time during peak hours?*
- **RQ7:** *Were there any significant delays or downtimes in processing transactions?*
- **RQ8:** *What were the best sale places?*
- **RQ9:** *What were the best top-up points?*
- **RQ10:** *Who were the best vendors?*
- **RQ11:** *What were the best products?*

The current platform already provides some performance metrics, but these questions should provide more detailed insights into the performance of the event.

Beverage Consumption

- **RQ12:** *What was the total amount of beverage consumed?*
- **RQ13:** *How many returnable cups were issued and returned/not returned?*
- **RQ14:** *What was the most popular beverage category?*

- **RQ15:** *What was the top beer brand, and how much was consumed and sold?*
- **RQ16:** *What was the most popular brand of other alcoholic beverages, and how much was consumed and sold?*
- **RQ17:** *What was the most popular non-alcoholic beverage brand, and how much was sold?*

Currently, a more in-depth product analysis is missing in the platform and the most important part of the product sales analysis at festivals is the beverage consumption. These questions should try to answer and give detailed insights into the beverage consumption, preferences and sales at the event.

Customers

What is the distribution of card schemes used to top up credit both on-site and online?

- **RQ18:** *What was the total attendance at the event, and how many active customers were there each day?*
- **RQ27:** *How many customers added credit in advance online?*
- **RQ28:** *What was the distribution of customers by type (on-site, online, staff, guest, VIP)?*
- **RQ29:** *How many customers use the mobile app?*
- **RQ19:** *What is the distribution of bank cards used to refund credit?*
- **RQ20:** *W*
- **RQ21:** *What was the course of the event in terms of new visitors? And when were the largest “peaks”?*
- **RQ22:** *What is the average time of a visitor from arrival to first transaction?*
- **RQ23:** *What was the course of the event in terms of topping up credit on-site? And when were the largest “peaks”?*

- ***RQ24:*** *How many customers topped up their credit multiple times, and how much only once?*
- ***RQ25:*** *What were the beverage preferences throughout the day?*
- ***RQ26:*** *What were the most common product combinations?*

The customer analysis is the most important part of the data analysis. It is crucial for the festival organizers to know their customer base and their behavior to optimize the event and make it more attractive for the customers.

Currently, no customer analysis, other than the customer ratings and list of customer chips, is available in the platform.

Answering these questions will possibly lead to the most valuable insights about the event's customer base and their behavior that no other platform or system currently provides.

Making these questions crucial and most valuable for the organizer and for the platform itself.

Technical Objectives

Answering the above questions will require a technical solution that will be able to process the data and provide the answers.

The scope of this study is not to implement a new system or a new platform and not even to implement any new changes to the existing NFCtron Hub platform. It is to find answers to the above questions and present them in a clear and understandable way in the form of a simple internal dashboard.

The technical goals of this study are:

- Prepare, process and analyze the data from **The Event**.
- Find answers to the above questions.
- Implement a simple internal dashboard that will present the answers to the questions.

Scope of the Study

To ensure the feasibility and focus of the study, certain boundaries have been defined in terms of what is included and excluded from the scope of the study.

Included in the Scope

- The study will focus on transactional, customer, and operational data from a specific event, referred to as **The Event**.
- Key areas of analysis include cashflow, revenue sources, performance indicators, beverage consumption, and customer segmentation and behavior.
- The data analyzed includes pre-event (e.g.,online top-ups), during-event (e.g.,chip transactions, sales), and post-event data (e.g.,credit refunds).
- A prototype dashboard will be developed using Python's Dash and Plotly libraries to present the key insights.
- The dashboard is intended for internal use and post-event analysis by the event organizer.

Excluded from the Scope

- **Real-Time Monitoring:** While the dashboard may be designed with real-time data potential, this study will focus solely on post-event analysis.
- **Multiple Event Comparisons:** This study is limited to the analysis of a single event (**The Event**) and does not involve comparative studies.
- **Data Collection:** The study does not involve the collection of new data and relies on data provided by **The Organizer** and the NFCtron system.
- **Implementation in NFCtron Hub:** The thesis focuses on analyzing data and developing a standalone prototype dashboard, not on direct integration into the NFCtron Hub platform.

Limitations

- **Anonymized Data:** To protect privacy, all customer and vendor data has been anonymized, which may restrict the analysis in some ways.
- **Single Event Focus:** Insights and recommendations are based only on data from **The Event**, which may limit broader generalizations.
- **Time Constraints:** Due to the thesis's timeline, advanced features (e.g., predictive analytics) and technical implementations were deprioritized, but will be considered for future work.

1 Data and Methodology

This chapter addresses the process and challenges of local environment setup, obtaining, preparing and anonymizing the data. Most importantly, this chapter describes and explains the data that was used for this research. It also briefly describes the tools, technologies and methods employed to answer the research questions.

1.1 Environment and local setup

To start off, we needed to set up some kind of environment where we would later work with the data. The data we would be working with was stored in a PostgreSQL database.

Having direct access to the production database to perform the analysis was not a secure and ethical way to go. Exporting only the necessary and raw data from the production database was an initial thought, but we initially did not know what data we would need, and by exporting we would lose all the relations between the tables.

Therefore, we decided to set up a local database with the same structure as the production database where we can query and analyze the data safely. The next step was to import the data from the production database to the local database. Importing or simple cloning the full database was also not an option, because only a small fraction of its subset was required.

So a deep internal analysis of the tables that were relevant to our study was performed. This resulted in a list of total 21 tables that held the necessary data for the study and were necessary to be imported.

1.1.1 Data Obtaining and Preparation

Almost every table was easily queried for the event and exported from the production database to a local CSV file. But some tables (for example and not surprisingly, the *transactions* table with over 140k rows) were too large to be exported in one piece, so we had to split them into smaller parts. Later, these parts were joined together to a single CSV file using a simple Python script.

Since no direct access to the production database was used for the export but rather a database management tool, the export was not as fast as it could be and took a significant amount of time. Moreover, the exported data, most importantly the timestamps, were in a different format than we needed. And also all numeric values were exported as a formatted strings with a comma as a decimal separator. So a data preprocessing Python script was written to convert such invalid columns to the correct format.

1.1.2 Local Database Setup

Then a step to set up the local PostgreSQL database was needed. Due to the nature of this study, we wanted to keep the setup as simple as possible, so we used the default PostgreSQL installation without using any special environment using Docker or similar. However, during this process I made a mistake and forgot that a PostgreSQL with PostGIS extension was needed. So it required to re-set up the database with the PostGIS extension.

The next step was to import the data from the CSV files to the local database. For further database handling, analysis and visualization, we used DataSpell, a Python IDE with a built-in database explorer and data visualization tools. DataSpell was then used for the local database import, which prior to it required some necessary database relations and constraints modifications, since the data was exported without them and was not relevant for the study.

This whole process resulted in approximately 387k rows of data in the local database that were ready to be queried and analyzed.

1.1.3 Local Database Modifications

Before any analysis was performed, some modifications to the local database were needed due to some known limitations and missing data.

Beverage Volumes

The first necessary limitation that the Beverage Consumption Analysis section heavily relied on was the missing information about beverage products volume in milliliters. This information was crucial for the analysis, so a new column was added to the relevant product information tables. However, the next step was to back-fill this information which was not easily automated.

The First approach was to write a Python script that would try to find the volume information from the product name. This worked for some products, but not for all since the naming convention was not consistent.

After several attempts to automate this process, it was decided to manually fill in the missing information since only 425 products were present in the database. Only 159 of them were of beverage type and thus eligible for the volume information.

Returnable Products

Since one of the research questions was to analyze the returnable cups and this information was not easily available in the database, a new column was added to the product information tables.

This was a simple binary column that indicated whether the product was returnable or not. Back-filling this information was also pretty straightforward since only one product was a returnable cup.

Venue Map Visualization

One of the initial ideas was to visualize the venue map with the locations of the selling places, top-up service points, stages and other important places.

This would be invaluable, the database was partially ready for this, but the data would be significantly time-consuming to back-fill and the later analysis and visualization would require more time.

Since these facts and the fact that this process of preparing the data took place before completing the list of data analysis questions, this idea was later abandoned.

Event Program

To present some time-related data and its correlation with the event program, it would require to have the event program in the database.

Again, the database was ready for this, but no event program was set up, since it was unnecessary for the event. Therefore, this required getting the event program from the festival website and manually insert it into the database.

This was manually a very time-consuming process, but it was necessary for the analysis. For some simplification of the process, an AI tool was used to extract the data from the program schedule screenshots and instructed to prepare an SQL script that would insert the data into the database.

This seemed like a good idea, but the AI tool was initially hallucinating and made up some incorrect data. But after several iterations, it successfully extracted the data and prepared the SQL script which was used and the event program was successfully inserted into the database.

In the end, I doubt that this process was faster than manual data entry, but it was a good exercise and a good example of how AI can be used to automate some processes.

1.2 Data Anonymization

The Data Anonymization process was necessary due to requirements initially set by the data provider and later by the ethical considerations. This step was performed for the already imported data in the local database. It required identifying the sensitive data and replacing them with anonymized values.

In this case, the most sensitive data were:

- **Vendor names:** Since it included the legal names of the vendors, it was necessary to anonymize them.
- **Selling places:** Some selling places were named after the vendors, so it was necessary to anonymize them as well.
- **Customer information:** Some tables included customer information like names, emails, phone numbers, etc.

The process could have been done various ways, but the fact that this study will not be exposing internal database structure, it was decided to perform the anonymization directly in the database.

However, if one-way anonymization were to be performed, it would permanently overwrite the original data, losing the possibility to switch from anonymized to original data. Therefore, a two-way anonymization process was chosen and performed.

This was particularly useful during the analysis phase, where the results would contain the original data for better understanding, fact-checking and for the internal presentation and consultations with the organizer.

It was done on the database level, where two new internal tables were introduced – *public.anonymization_config* and *public.original_values*.

Where the *public.anonymization_config* table held the configuration about which schema, table and column should be anonymized and how. For the usage, a simple SQL function was created to define the anonymization configuration in a simple JSON format, that looked like in the **Anonymization configuration example..**

```
1  SELECT configure_anonymization('[
2    { "table": "schema.seller", "columns": ["legal_name", "name_int",
      ↳ "name_pub"] },
3    ...
4    { "table": "schema.user_account", "columns": ["email", "last_name",
      ↳ "first_name", "phone"] },
5    ]'::JSONB
6  );
```

Source code 1: Anonymization configuration example.

The *public.original_values* table was used to store the original values of the anonymized columns. Again, using a simple SQL function *anonymize_database()*, it would store the original values and anonymize the configured table columns.

One particular challenge was anonymizing the values smartly. It could have been easily done by replacing the values with random strings, hashes or encrypted values. But working with data, where a vendor is named *fa65165b923e9cc* is not very convenient.

Therefore, a simple SQL function was written to anonymize the value depending on the configuration. This allowed to configure the anonymization to:

- replace vendor names with values like *Vendor 1*,
- customer emails with *03b09592-d0eb-43a3-9941-30d38ade6bce@gmail.com* keeping the original email domain,
- selling places with values like *Place 1* where the original name contained sensitive information, but keep original values for places like *BAR L2*, etc.

In the end, it resulted in a database with anonymized values per stated configuration, that could have been used for the analysis and results presentation safely. However, the original values were still present in the database, and the database could have been anytime easily restored to the original state if needed and vice versa.

1.3 Data Structure

Without exposing the internal database structure, the abstract data structure this study was working with can be described as follows:

1.3.1 Financial Data

Transactions (approx. 300k rows): Transactional data, that holds the information about the type, amount, timestamps and links to products, places and other related entities. This analysis relies on and works with several transaction types,

including **top-up charge** and **refund** transactions ¹, **order sales** and **refund** transactions ² and **chip registrations**³.

Credit Refunds (approx. 15k rows): Post-event credit refund requested by the customers with the information about the amount, timestamp, customer and the bank account to which the refund was sent. This data enables our analysis to correctly handle disposable credit balances, more information about the anonymous customers ⁴ and their behavior.

Why is it important?

These records also form the backbone of the analysis, enabling insights into revenue, sales trends, and customer behavior.

1.3.2 Customer Data

User Accounts (approx. 5k rows): Registered customer accounts with the information about the user and its potential online order history and other related information.

Tickets and Orders (approx. 30k rows): Information about the online sold tickets, its types, prices, timestamps, online order related information.

Why is it important?

With the above data, this analysis can work with more customer information supporting the customer behavior analysis, customer segmentation and other related analysis.

1.3.3 Event Data

Places (approx. 400 rows): Selling and top-up service points, zones for access control and its other relations.

Products (approx. 500 rows): Essentially a product catalog including the product name, price, category, volume, seller ownership and seller-organizer deal re-

¹Top-up transactions mean funding or refunding chip credit balances.

²Order transactions mean spending the credit balance for products.

³Chip registrations mean records of when the chip bracelets were registered by the system.

⁴Anonymous customer essentially means a user without registered account.

lated links.

Important information about products is their supported categorization that will later be used for the sales analysis and can be seen in Table 1.1.

Category	Description
Nonalcoholic	Any non-alcoholic beverages (e.g., coffee, water, etc.)
Beer	Any kind of beer.
Wine	Any kind of wine.
Other Alcohol	Any other kind of alcoholic beverages (e.g., shots, cocktails, etc.).
Salty	Any salty snacks.
Sweet	Any sweet snacks.
Other	Any other products that do not fit into the above categories.

Table 1.1: Product categories

Event Program (approx. 140 rows): Event program schedule with the information about the stages, performers, times and other related information.

Why is it important?

This data provides more context to the event when combined with the financial and customer data above. Enabling the analysis to work with the event program, its correlation with the sales, customer behavior and other related analysis.

1.3.4 Data Processing Views

To efficiently query the studied data during the analysis, several SQL views and functions were created to simplify and speed up the process.

Transaction Commission Calculation: A function that calculates the commission for each transaction based on the product and the seller-organizer deal. This was a crucial method required to calculate and analyze the commission from event order sales contributing to the organizer's revenue.

Transaction Enrichment: Since the transactional data consisted of several transaction types which were not easily distinguishable, a view was created to enrich the transaction data with the transaction type information. It also benefited from the transaction commission calculation function mentioned above which helped to easily calculate the commission for each transaction.

Chip Customers: Probably the most complex function that returns the customers at the event. Since the transactional data is architected using Event Sourcing, the customer information is not directly available and needs to be compiled from the transactional history. This function was constructed in a way, where it supports time-based filtering and provides extensive insights into the customers, which is shown in Table 1.2 below:

Column Name	Description
CHIP_ID	Unique chip identifier.
CHIP_TYPE	Type of chip (e.g., regular, VIP, online, staff).
REG_AT	Timestamp of chip registration.
FIRST_TRX	First transaction timestamp associated with the chip.
LAST_TRX	Last transaction timestamp associated with the chip.
LAST_BALANCE	Last known balance at the specified time frame.
ACTUAL_BALANCE	Balance at the specified time frame after credit refunds.
IS_BLOCKED	Indicates if the chip is blocked due to suspicious activity.
HOURS_ACTIVE	Total active hours of the chip (daily sum).
DAYS_ACTIVE	Total number of days the chip was active.
T_COUNT	Total number of transactions associated with the chip.
O_TOTAL_CNT	Total number of orders placed using the chip.
O_TOTAL_AMT	Total amount spent through orders.
O_MAX_AMT	Maximum amount spent in a single order.
O_AVG_AMT	Average amount spent per order.
O_MODE_AMT	Most common amount spent per order.
OS_AVG_AMT	Average amount spent on sales orders (excluding refunds).
OS_MODE_AMT	Most common sale amount (excluding refunds).
TU_TOTAL_CNT	Total number of top-ups made to the chip.
TU_TOTAL_AMT	Total amount credited to the chip via top-ups.
TU_MAX_AMT	Maximum amount credited in a single top-up.
TU_AVG_AMT	Average amount credited per top-up.
TU_MODE_AMT	Most common amount credited per top-up.
TU_CARD_BRAND	Used card brand for top-ups (e.g., Visa, Mastercard).
BR_AMT	Total amount refunded to the customer's bank account.
BR_EMAIL_DOMAIN	Domain of the refund request email (e.g., gmail.com).
BR_COUNTRY	Country associated with the bank account for refund.
BR_REQ_SOURCE	Source of the refund request (e.g., iOS, Android, Web).
BR_BANK_NAME	Name of the Czech bank used for the refund.
BR_CREATED	Timestamp when the refund request was created.
BR_APPROVED	Timestamp when the refund request was approved.
A_EMAIL_DOMAIN	Email domain of the account.
A_COUNTRY_NAME	Country associated with the account.
A_REQ_SOURCE	Source of the account creation (e.g., iOS, Android, Web).
EO_PAYMENT_METHOD	Payment method for orders (e.g., card, bank transfer).
EO_CARD_BRAND	Used card brand for online order (Visa, Mastercard)
EO_REQ_SOURCE	Source of the order request (e.g., iOS, Android, Web).

Table 1.2: Customer chips function return table.

For further understanding, an enumeration of chip types should be mentioned; that can be seen in Table 1.3.

Chip Type	Description
Regular	Chip issued at the event without any prior credit.
VIP	VIP chip issued with artificial credit on behalf of the organizer.
Online	Chip issued via online purchase or top-up
Staff	Chip issued to the staff members.
Guest	Chip issued to the guests (band members, etc.)

Table 1.3: Chip types enumeration

1.4 Tools and Technologies

As mentioned earlier, this process utilized a variety of tools and technologies to handle the data and prepare for the analysis. These main tools and technologies included:

- **PostgreSQL:** An open-source relational database management system used to store and query the data.
- **PostGIS:** An extension of PostgreSQL that supports geospatial data, enabling spatial analysis and visualization.
- **DataSpell:** An integrated development environment (IDE) for data science and analytics, used for database exploration and data visualization.
- **Python:** Used for data preprocessing, querying, and analysis, along with libraries like Pandas for data manipulation.
- **Claude AI:** Utilized for extracting data from unstructured sources (e.g., program schedules) and automating repetitive tasks like data entry.

Using such tools during this process provided a convenient environment for data handling and processing, initial analysis and ensuring the accuracy and efficiency for further analysis and results presentation.

1.5 Conclusion

This chapter laid out the process of setting up the local environment, collecting and preparing the data, anonymizing sensitive information and introducing the data structure.

Some key challenges, such as missing beverage volume information, data anonymization, and event program integration, were addressed through a combination of manual and automated processes.

In the end, the use of SQL views for data enrichment and functions for further data processing ensured efficient querying and analysis. And hopefully laid the groundwork for answering the research questions and achieving the study's goals, as detailed in the later chapters.

For better understanding and visualization of this initial comprehensive Knowledge Data Discover process, a simplified diagram was created; that can be seen in Figure 1.1.

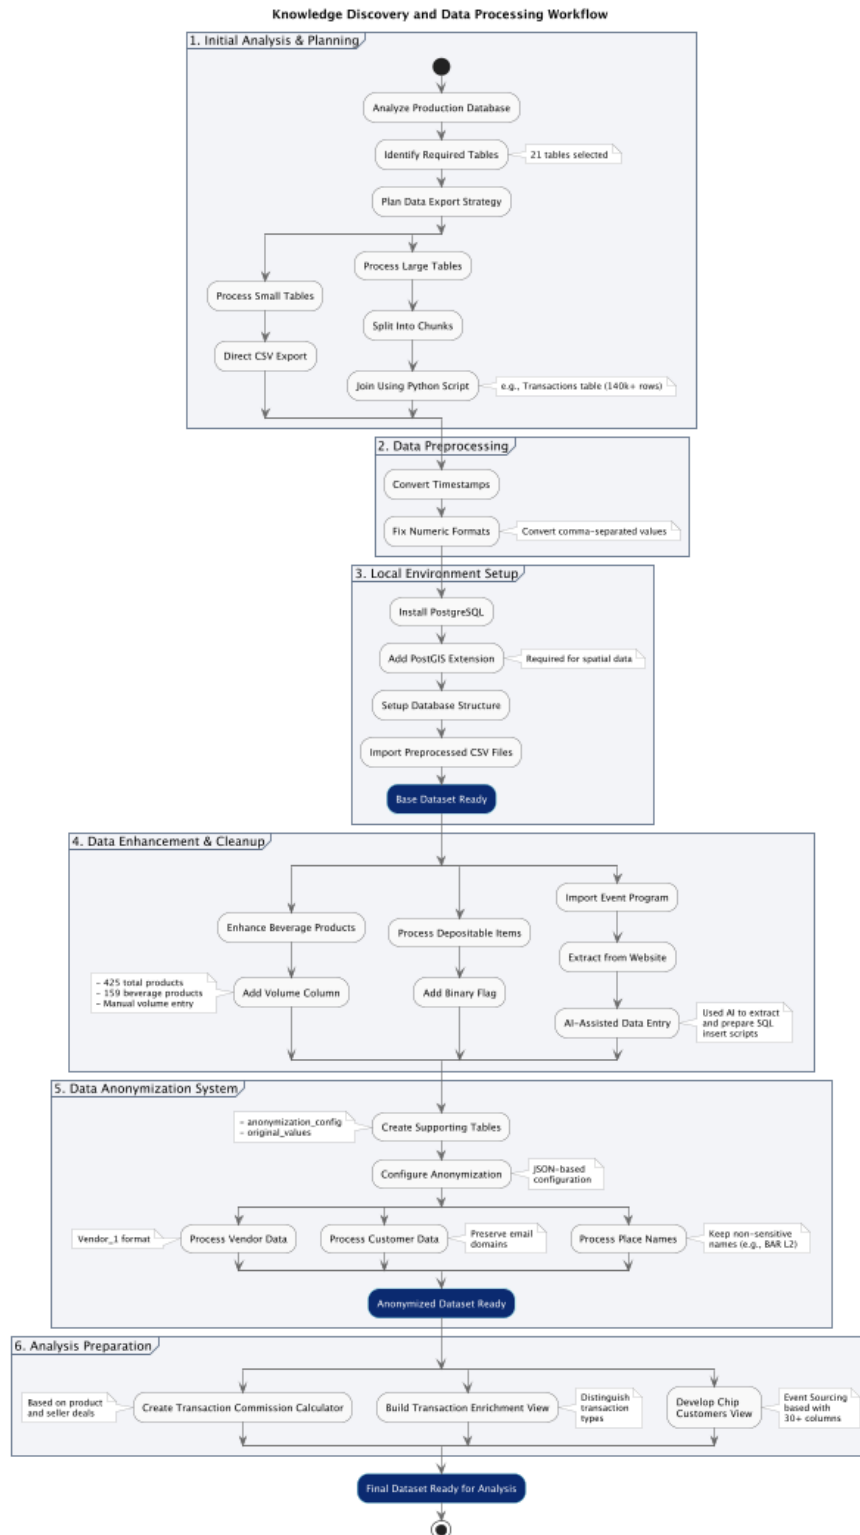


Figure 1.1: Knowledge Data Discovery workflow diagram.

2 Data Analysis and Results

This chapter presents the data analysis and results of the research. The goal is to address the research questions outlined earlier, with a focus on providing actionable insights for the event organizer.

The chapter is divided into several sections corresponding to key analytical areas:

1. **Cashflow and Revenue Sources Analysis,**
2. **Performance Indicators Analysis,**
3. **Beverage Consumption Analysis,**
4. **and Customer Analysis**

Each section focuses on a different aspect of the data analysis trying to answer the research questions, present quantitative results, visualizations, and interpretations.

2.1 Cashflow and Revenue Sources Analysis

This section provides a comprehensive view of the festival's financial performance and cash flows. It should answer critical questions about how finances were funded into the system, how were they processed, and what were the final outcomes.

For this analysis, four questions were previously formulated. However, they were reordered to better fit the narrative of the analysis and logical flow of the chapter:

- ***RQ2:*** *How much and how was the balance topped up on the chips?*

- **RQ4:** *What was the total sales of the event, how much of it was sold by the organizer, and how many external vendors?*
- **RQ3:** *What is the remaining balance on all chips after the event and refunds?*
- **RQ1:** *W*

In the end, this section should provide a clear picture of the financial flows during the event and easy understanding of the generated revenue from various sources.

2.1.1 Chip Top-Up Analysis

Research Question


RQ2: How much and how was the balance topped up on the chips?

Attendees could top up their chip balances via online prepayments or on-site using cash or card. Additionally, the system allows to top-up “artificial” credit for VIP-issued chips which is also a mean of funding the system. However, these VIP credits are later not refundable, but this will be discussed in the next section.

This subsection quantifies these methods, highlighting their respective contributions to the overall top-up total.

To get the results, it was necessary to find all top-up transactions and their respective payment methods used. This resulted in **17,704** top-up transactions, with a total value of **14,520,973 CZK**.

When looking at the grouping by payment methods, the results in Figure 2.1 give a clear picture of the distribution.



Payment Method	Count	Total Value (CZK)
Card terminal	8,486	7,264,503 CZK
Cash	7,561	5,782,570 CZK
Online pre top-up	1,634	1,436,400 CZK
VIP issued	23	37,500 CZK

Figure 2.1: Top-Up Transactions by Payment Method

Source: *Own creation*

Thanks to the results, it is clear how many funds did the system receive and by what means.

Key Takeaways

- Total top-up amount was **14,520,973 CZK**.
- Most used payment method was card terminal at the event with 50% of all top-ups.
- Only around 10% of the top-ups were done online.

2.1.2 Sales Analysis

Research Question

RQ4: What was the total sales of the event, how much of it was sold by the organizer, and how many external vendors?

The sales analysis was crucial for understanding the overall sales behavior and served as a basis for further insights tightly connected to the revenue sources.

To answer the research question, it was necessary to find all sales transactions and their respective sellers and to divide them into two groups: the direct organizer's sales and external vendors' sales. And for better understanding, the sales were also grouped by the product categories (see Table 1.1 for the list of categories).

The results show that the total sales of the event were **11,711,807 CZK** with the organizer's sales being **8,240,264 CZK** and the external vendors' sales **3,471,543 CZK**.

The organizer, most importantly, sold all the beer beverages and most of the non-alcoholic and alcoholic (spirits) beverages. Whereas the external vendors sold mainly the food, wine beverages and other uncategorized products. This can be seen in Figure 2.2 below.




Figure 2.2: Sales of the Organizer vs. External Vendors

Source: *Own creation*

The organizer also sold not so little of uncategorized products, which after further investigation turned out to be ticket sales at the event amounting to **684,700 CZK**.

In total, the organizer direct sales were **70%** of the total sales, which is a signifi-

cant portion, and thus the organizer itself has even bigger influence on the event's financial performance.

Key Takeaways

- Total sales of the event were **11,711,807 CZK**, where organizer sales were **70%** of the total.
- The organizer sold all beer beverages and the majority of the non-alcoholic and alcoholic beverages.
- The organizer also sold tickets at the event amounting to **684,700 CZK**.
- External vendors sold mainly food, wine beverages, and other uncategorized products.

TODO: Better chart

2.1.3 Remaining Chip Balances

Research Question

RQ3: What is the remaining balance on all chips after the event and refunds?

The remaining chip balances are crucial for the event organizer as they represent the potential revenue that can be still claimed. Any unclaimed balances after a given refund period, which is usually up to 14 days after the event will be considered as organizer's taxable revenue.

Out of the total top-up amount of **14,520,973 CZK**, the total spent credit amounted to **10,984,945 CZK**, which left a total of **3,536,028 CZK** on the chips before refunds. After refunds – done both at the event (**15,379 CZK**) and later via online bank refund requests (**3,163,567 CZK**) – the remaining balance was reduced to **357,082 CZK**.

However, this still included the artificially issued VIP credits with leftover balance of **12,405 CZK**. The system also reported integrity errors in the data, which resulted in a total of **10,246 CZK** due to fraudulent activities performed by some attendees which were automatically suspended by the system.

This left the total unclaimed balance at **334,431 CZK**, which has been claimed by the organizer as taxable revenue.

Since these numbers can be quite abstract, the results in a form of sankey diagram in Figure 2.3 below provide a clear picture of the flow of the funds.

TODO: Better chart




Figure 2.3: Remaining Chip Balances Sankey Diagram

Source: *Own creation*

Thanks to this breakdown, it is clear how the remaining balances were reduced and what was the final outcome. These results are important for the last part of this section, which is the total revenue of the organizer.

Key Takeaways

- Total unused credit was **3,536,028 CZK**.
- Credit refunded to customers was **3,178,946 CZK**.
- After VIP issued credits and system integrity error, the unclaimed balance was **334,431 CZK**.

2.1.4 Total Revenue of the Organizer

Research Question

RQ1: W

The festival's financial model is based on a combination of revenue streams.

The most important stream is the **commission from the vendor sales**, which is arranged in advance between the organizer and the vendors. The commission is, in this case, a percentage (ranging from 15% to 30% depending on the deal) of the vendor sales amount without VAT.

Therefore, this required finding all sales transactions made at the external vendors' stands and calculating the commission based on the agreed percentage. However, this was not a straightforward task, since a transaction could contain multiple products even from different vendors.

This required a more complex calculation, for which was used the previously mentioned data processing views which were designed for this purpose. In the end, the total revenue from sales commissions was **820,712.79 CZK**.

Another source of revenue is the **unclaimed chip balances**, which, after a credit refund period, are considered as taxable revenue for the organizer. This, thanks to the previous subsection, was found to be **334,431 CZK**.

TODO: Better chart

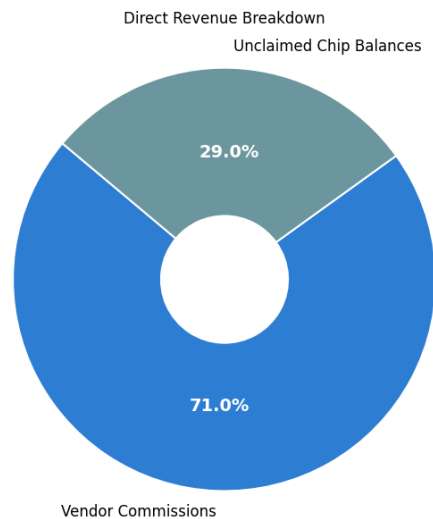


Figure 2.4: Breakdown of Direct Revenue Streams

Source: *Own creation*

Currently totalling **1,155,143.79 CZK** is the direct revenue of the organizer from the event and can be seen in Figure 2.4 above.

However, given the circumstances and setup of this event, there were also additional, but indirect revenue streams that were not included in the total revenue. These include **the online ticket sales**, which were sold by the organizer and **the direct sales of the organizer**. They were not included in the total direct revenue, as they may misinterpret the results since the analysis lacks expenses of the organizer.

If we were to include these, the total revenue would increase by **11,179,700 CZK** from the online ticket sales and **8,240,264 CZK** from the direct sales, which would result in a total revenue of **20,575,107.79 CZK**.

To better understand the revenue streams, the results are visualized in Table 2.1 and in Figure 2.5 below.

Revenue Stream	Amount (CZK)
Vendor Commissions	820,712.79 CZK
Unclaimed Chip Balances	334,431 CZK
Total Direct Revenue	1,155,143.79 CZK
Online Ticket Sales	11,179,700 CZK
Organizer Direct Sales	8,240,264 CZK
Total Revenue (All Streams)	20,575,107.79 CZK

Table 2.1: Revenue Summary Breakdown

TODO: Better chart




Figure 2.5: Breakdown of All Revenue Streams

Source: *Own creation*

Key Takeaways

- Total direct revenue of the organizer was **1,155,143.79 CZK**.
- Vendor sale commission contributed to approximately 71% of the total direct revenue.
- With other indirect revenue streams, the total revenue would be **20,575,107.79 CZK**.

2.1.5 Summary

This section provided a comprehensive view of the festival's financial performance and cash flows. The results covered the top-up transactions, sales analysis, remaining chip balances, and the total revenue of the organizer and contributed to a better understanding from the financial perspective of the festival.

Nevertheless, results covered in these subsections are only a part of the whole picture and can be interpreted in various ways.

For this particular challenge, a summarized cash flow diagram of payments was created, containing thus only the direct revenue streams. This diagram can be seen in the Figure 2.6 below.

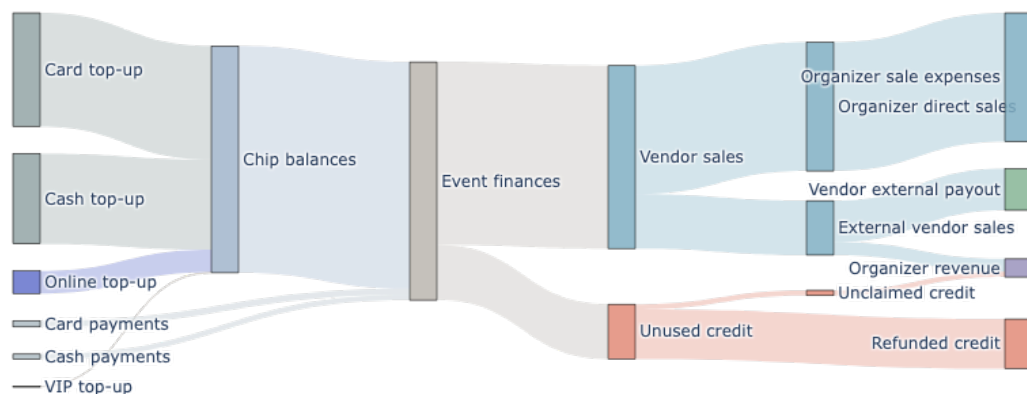


Figure 2.6: Overall Cash Flow Diagram

Source: *Own creation*

This diagram provides a clear overview of the financial flows during the festival and nicely summarizes the results of this analysis.

Key Takeaways

- Total incoming money flow was **14,520,973 CZK** from top-up transactions and **726,862 CZK** from non-chip sales.
- Total sales amounted to **11,711,807 CZK**.
- Which left a total of **3,536,028 CZK** in unused credit before refunds.
- After refunds and non-refundable chips, the remaining balance left was **334,431 CZK** claimed as taxable revenue.
- Commission from external vendor sales contributed to **820,712.79 CZK** of the total direct revenue.
- Together, the total direct revenue of the organizer was **1,155,143.79 CZK**.

2.2 Performance Indicators Analysis

This section focuses on the performance indicators of the event. The goal is to identify key metrics that can be used to further evaluate the event and its success. The potential of this analysis is to measure the “greatness” and the size of the event in terms of performance.

This analysis sets out to answer other previously defined research questions related to the performance of the event. Which, for purposes of this analysis, were slightly reordered and grouped into the two following subsections:

1. Transactions Processing Analysis

- **RQ5:** *How many transactions were processed in total, and when was the system’s largest “peak” in transaction volume?*
- **RQ6:** *What was the average transaction processing time during peak hours?*
- **RQ7:** *Were there any significant delays or downtimes in processing transactions?*

2. and Best Sale & Top-Up Points, Vendors, and Products Analysis

- **RQ8:** *What were the best sale places?*
- **RQ9:** *What were the best top-up points?*

- **RQ10:** *Who were the best vendors?*
- **RQ11:** *What were the best products?*

The results of this analysis should provide insights into the event’s performance and help the organizer to understand the key metrics that can be used to evaluate the event’s success.

2.2.1 Transactions Processing Analysis

This subsection will focus on the processing of transactions during the event in pursuit of answering the three first research questions of this section.

Research Question

RQ5: How many transactions were processed in total, and when was the system’s largest “peak” in transaction volume?

This question actually consists of two sub-questions, which will be addressed separately.

The first part questions the total number of transactions processed during the event. Which was actually pretty straightforward to answer, as the system was designed to track all transactions and their respective types. The resulted total number of transactions was **141,381** consisting of **110,854** sales transactions, **17,726** top-up transactions and **12,801** chip register transactions.

The second part focuses on rather time-related metrics and asks about the processing peak times during the festival. For this part, it was necessary to spread out the above transactions over the time and find the peaks.

The results in Figure 2.7 below show the distribution of the processed transactions over time. It clearly identifies the peak on the last day of the festival at 18:00 amounting to **8,986** transactions.

TODO: Better chart




Figure 2.7: Transactions Processing Peaks

Source: *Own creation*

Key Takeaways

- Total number of transactions processed was **141,381** consisting mainly (78%) of order transactions.
- The peak of transactions was on the last day of the festival at 18:00 with **8,986** transactions processed at that hour.

The following two questions (RQ6 and RQ7), focus on processing times and potential delays during the event.

Research Question

RQ6: What was the average transaction processing time during peak hours?

The answer to this question is closely related to the previous one, as it requires the identification of the processing times during the peak times, which were already identified.

It required finding the average processing time, meaning difference between the transaction creation and its completion times.

? What causes the processing time?

Time when the transaction is created is the time when the in-place offline-supported system created the transaction, and the processed time is later when the central system receives the transaction and processes it. The delays can be caused by various factors, such as network latency, offline mode active, system load, or even the transaction type.

The results show that the average processing time during the peak times was approximately **40** seconds.

When slightly changing the displayed data, we also get the answer to the RQ7 about the potential delays and downtimes during the event.

Research Question

RQ7: Were there any significant delays or downtimes in processing transactions?

The chart in the Figure 2.8 shows the distribution of the processing times over the time and identifies one high processing peak of approximately **13** minutes. This is highly unusual and indicates a vendor misuse of the system or accidentally put the system into offline mode.

TODO: Better chart




Figure 2.8: Transaction Processing Peaks

Source: *Own creation*

Other high peaks are visible on the second day in the afternoon with **5** minutes

processing time, which was probably caused by the initial load on that day.

Key Takeaways

- The average processing time during the peak times was **40** seconds.
- The highest processing peak was approximately **13** minutes, indicating a potential misuse of the system.
- Other high peaks were visible on the second day in the afternoon with **5** minutes processing time.

These results provide insights into the system's performance during the event, its reliability, and potential bottlenecks. It also shows the festival's popularity and the system's ability to handle the load.

2.2.2 Best Sale & Top-Up Points, Vendors, and Products Analysis

In this subsection, the main goal will be to address the last four research questions of this section and provide insights into the best: selling points, top-up points, vendors, and products.

The problem with these question statements is that they are quite broad and can be interpreted in various ways. What does a “best” mean in this context?

It can be the most profitable, the best rated, the most visited, etc. But since we are exploring the performance indicators, the best should be understood as the “busiest”. Which in terms of the system and this analysis should mean **the most transactions created** and the point's ability to handle the load.

Best Top-Up Points

The first focus will be on the best top-up points since unlike the selling points, vendors and products, the top-up points are not linked to any specific product or vendor.

Research Question

RQ9: What were the best top-up points?

To find these results, it required finding all top-up transactions, aggregate them in a bucket-like time frame and finally calculate their total counts, max peaks and averages over time.

This resulted in the following findings in the Table 2.2 below.

	Top-Up point	Customers	Transactions	Max trx./h
1	Pokladna 16	1,119	1,139	99
2	Pokladna 2	1,103	1,108	106
3	Pokladna 15	1,035	1,043	82
4	Pokladna 4	1,007	1,017	107
...				
15	Pokladna 7	740	744	73
16	Pokladna 9	699	711	72
17	Pokladna 10	629	640	69
18	Odbavení	529	529	125
...				
25	Odbavení 5	198	198	46
26	Odbavení 6	191	191	32
27	Support	64	85	11
28	VIP	23	23	5
29	Akreditace	17	17	8

Table 2.2: Best Top-Up Points

This indicates that the most busy top-up points were somehow evenly distributed with approximately around **1000 transactions** processed during the event with average peaks of around **100 transactions/hour**. The least busy top-up points were the specific ones, such as the Support tent, VIP and Accreditation points, which were not used so much for top-ups.

The overall distribution, shown in the Figure 2.9, also shows that Top-up points were more busy than Check-in points. That makes sense because the top-ups were done more frequently than the initial check-ins, but the check-ins were done in a more concentrated time frame.




Figure 2.9: Best Top-Up Points

Source: *Own creation*

TODO: Better chart

Especially **Odbavení** point processed only **529** transactions but peaked at **125** transactions per hour, which was much higher than the other check-in points and even higher than the best top-up points.

Key Takeaways

- The most busy top-up points processed around **1,000** transactions during the event.
- The average peak of the top-up points was around **100** transactions per hour.
- The least busy top-up points were the specific ones, such as the Support tent, VIP and Accreditation points.
- Check-in points were less busy than the top-up points, but the **Odbavení** point peaked at **125** transactions per hour.

Best Sale Points

The next focus will be on the best sale points, which are the points where the most orders were created.

Research Question

RQ8: What were the best sale places?

The process of finding the best sale points was similar to the previous one, but this time it required finding all sales transactions and their respective points.

Out of total of **145** sale points, the best place was undeniable the **L20 PIVNÍ STAN 1** with total **10,114** orders processed and the maximum peak of **840** orders per hour.

Another interesting fact is the number of unique users processed at the places. The **L20 PIVNÍ STAN 1** processed **9,159** unique customers, accounting for a significant portion (**91.50%**) of the total.

Based on this particular finding, we can assume that in the following analysis - the best vendors and products - the product preferences will be highly in favor

	Sale Point	Customers	Orders	Max orders./h
1	L20 PIVNÍ STAN 1	9,159	10,114	840
2	Place 78	4,914	5,113	479
3	A19 PIVO 1	4,615	5,073	608
4	L16 FRISCO	3,420	3,646	235
5	Place 11	3,078	3,226	265
6	L22 BEEFEATER/ HAVANA 1	2,613	2,857	317
7	A19 VÝKUP KELÍMKŮ	2,683	2,743	316
...				
132	Support	23	23	4
133	Taxi 51	11	21	4
134	Taxi 50	5	8	4

Table 2.3: Best Sale Points

of the beer beverages. And thus the best vendor will probably be the organizer, as they sold all the beer beverages at the festival.

Key Takeaways

- The best sale point was the **L20 PIVNÍ STAN 1** with total **10,114** orders processed, which was **9.12%** of the total orders created.
- The best sale point peaked at **840** orders per hour.
- The **L20 PIVNÍ STAN 1** processed **9,159** unique customers, which was **91.50%** of all active customers.

In conclusion to these two questions, the results show clearly the busiest points of the festival and their ability to handle the load. However, the results can be visualized in a more interactive way, which would provide a better understanding of the data.

One especially interesting visualization of the best sale and top-up points would be a heatmap of the festival area with the points and their respective transaction counts. As this was initially intended to be part of the analysis, it was unfortunately not possible to create it due to the lack of the necessary data.

Best Vendors

To analyze the best vendors, in terms of performance, the same approach as with the sale points will be used.

Research Question

RQ10: Who were the best vendors?

The results in the Table 2.4 below show the distribution of the processed orders over the vendors.

	Vendor	Customers	Orders
1	Organizer	9,831	89,217
2	Seller 05	3,632	5,113
3	Seller 15	1,571	2,282
4	Seller 13	1,533	1,812
5	Seller 29	915	1,623
6	Seller 28	530	1,468
7	Seller 08	1,118	1,359
...			
25	Seller 02	30	31
26	Seller 09	26	27
27	Seller 10	4	5

Table 2.4: Best Vendors

As predicted in the previous section, the best vendor was the organizer, which processed the most orders and customers. The total of **89,217** orders was processed by the organizer, which was **80.04%** of the total orders created and served **9,831** unique customers, which was **98.22%** of all active customers.

The second-best vendor, out of **27** total, was some **Seller 05** with about **84,104** orders less than the organizer.

In these results, we did not go after the hourly maximums, as in the previous questions, as the vendors were not time-bound and could have been at multiple places simultaneously. The results would not be so relevant and would not provide any additional insights.

Key Takeaways

- The best vendor was the organizer, which processed **89,217** orders (**80.04%** of the total orders).
- The best vendor served **9,831** unique customers (**98.22%** of all active customers).
- The second-best vendor was behind around **84,104** orders less than the best vendor.

Best Products

The last focus will be on the best products, which are the products that were sold the most during the event.

Research Question

RQ11: What were the best products?

Previously, the best products were predicted to be the beer beverages, as the organizer sold all the beer beverages at the festival.

The results in the Table 2.5 below somehow confirm this prediction as the very best product was a returnable cup – **Kelímek – záloha** and the next four best products were beer beverages:

	Product	Sales	Refunds	Customers
1	Kelímek - záloha	17,933	11,215	8,729
2	Radegast 12°	12,844	26	3,760
3	Hladinka - Prazdroj	9,106	21	2,958
4	Radegast 10°	7,662	12	2,846
6	Birell	4,367	6	2,321
...				
326	šnúřka s karabinou	1	0	1
327	Ponožky	1	0	1
328	Klobásky 2ks	1	0	1

Table 2.5: Best Products

As there were more than **300** products sold during the event, a better visualization of the results would be a bar chart showing the best product categories instead of individual products. This can be seen in the Figure 2.10 below.




Figure 2.10: Best Products by Category

Source: *Own creation*

TODO: Better chart

This chart now confirms the prediction about the beer beverages, as the **Beer** (Pivo) category was the most sold during the event with a little more than **40,000** orders processed.

Key Takeaways

- The best product was a returnable cup followed by beer beverages.
- Prediction about the beer beverages was confirmed, as the **Beer** (Pivo) category was the most sold during the event.

This last analysis provided insights into the best products, which confirmed the previous predictions and will now serve as a basis for the next section, where the beverage consumption will be analyzed.

2.2.3 Summary

Thanks to this section, the performance indicators of the event were analyzed, and the key metrics were identified. The results analyzed the transactional processing performance, identified several peak points during the event, and provided insights into the best sale points, top-up points, vendors, and products. Provid-

ing a better understanding of the event’s performance and giving more context for the next analysis dealing with the beverage consumption.

2.3 Beverage Consumption Analysis

This section will provide a detailed analysis of beverage consumption, which was previously established as the most significant aspect of the event.

It should offer insights into overall consumption, detailed information on returnable cups, the most often drank beverages, and the leading beverage brands, addressing the previously established inquiries.

This section will address the next six research questions, after a minor logical reorganization into three groups:

1. Returnable Cups Analysis

- **RQ13:** *How many returnable cups were issued and returned/not returned?*

2. Total Consumption Analysis

- **RQ12:** *What was the total amount of beverage consumed?*
- **RQ14:** *What was the most popular beverage category?*

3. and Popular Brands Analysis.

- **RQ15:** *What was the top beer brand, and how much was consumed and sold?*
- **RQ17:** *What was the most popular non-alcoholic beverage brand, and how much was sold?*
- **RQ16:** *What was the most popular brand of other alcoholic beverages, and how much was consumed and sold?*

2.3.1 Returnable Cups Analysis

Due to the little alteration of the local database and the previously referenced data in subsection 1.1.3, it became possible to monitor the returnable cups and their associated transactions.

This capability, previously absent, should enhance comprehension of product sales and the utilization of returnable cups.

Research Question

RQ13: How many returnable cups were issued and returned/not returned?

To figure out the results, it was essential to analyze the actual contents of the transactions rather than merely identifying transactions including returnable cups, as a single transaction may encompass many products and hence multiple returnable cups. The organizer sold the cups for a price of 70 CZK with a 21% VAT.

Upon calculating the total number of issued and returned cups, the results depicted in Figure 2.11 below illustrate the distribution of returnable cups throughout the event.




Figure 2.11: Returnable Cups

Source: *Own creation*

The data indicate that a total of **20,045** cups was distributed during the event, with only **17,322** cups returned, yielding a return rate of **78.60%**.

This, however, does not imply that the remaining **4,723** cups were lost or regarded as a loss, as the cups were purchased. Most unreturned cups can be presumed to have been retained as souvenirs.

Key Takeaways

- The total of **20,045** cups were issued during the festival.
- Only **17,322** cups returned resulted in a **78.60%** return rate.

2.3.2 Total Consumption Analysis

This focuses on the overall beverage consumption during the event. Again, thanks to local database modifications, it was possible to track the beverage consumption easily, as each product now had a volume attribute in milliliters.

Research Question

RQ12: What was the total amount of beverage consumed?

To find the results was quite straightforward, as it only required summing up the volumes of all sold products.

However, to present the results, it was convenient to group the products into categories and show the total consumption of each category and thus answering also the next question RQ14.

Research Question

RQ14: What was the most popular beverage category?

These results are shown in the Table 2.6 below and show a total of **29,342** liters of beverages consumed during the event with the beer category being the most consumed.

Beverage Category	Volume	Ratio
■ Beer	19,797 l	67.46 %
■ Non-alcoholic	6,987 l	23.81 %
■ Spirits / other alcohol	1,992 l	6.78 %
■ Wine	575 l	1.95 %
Total volume	29,342 l	100 %

Table 2.6: Total Beverage Consumption

These results serve as a basis for the next questions, which focuses on the most consumed beverage brands rather than product categories.

Key Takeaways

- The total of **29,342** liters of beverages were consumed during the event.
- The beer category was the most consumed with **19,797** liters consumed.

2.3.3 Popular Brands Analysis

This section explores beverage preferences, concentrating on the most popular brands within the leading categories:

- **Beer Brands Analysis,**
- **Non-alcoholic Brands Analysis,**
- and **Alcoholic Brands Analysis.**

Answering these issues required the identification and categorization of all beverage products, followed by the computation of their overall consumption. However, this was not entirely easy, as the products were not uniformly labeled. This indicated that beers labeled as “Radegast 10” and “Radegast 12” were not categorized together, resulting in biased outcomes.

To address this issue, the products ought to be categorized by their brand rather than by their name. This methodology seemed logical; nevertheless, the data lacked brand information, containing simply the product name.

A systematic approach would be to extend the database with brands and back-fill the products with a link to the brand. This would require a significant amount of work and time, which was not available at the time of this analysis.

A more straightforward and manual method was used, wherein the product’s “brand” was identified by extracting the essential parts of the product name, eliminating redundant elements such as volume, beer grade, or other details. This helped produce better results, however, still lack complete accuracy.

Beer Brands Analysis

Starting with the most consumed and most popular category – beer.

Research Question

RQ15: What was the top beer brand, and how much was consumed and sold?

The results in Figure 2.12 below illustrate the distribution of the most frequently consumed beer brands during the event.




Figure 2.12: Most Consumed Beer Brands

Source: *Own creation*

This distribution indicates that the most consumed beer brand was **Radegast**, with a total consumption of **10,366** liters and **27,329** units sold.

The second was **Pilsner Urquell**, which had half the consumption of **Radegast**, followed by the **Frisco** and **Proud** brands. A complete list of all ten brands from the festival is shown in the Table 2.7 below.

	Beer brand	Total consumption	Units sold
1	Radegast	10,366 l	27,329
2	Pilsner Urquell	5,258 l	14,137
3	Frisco	2,658 l	8,209
4	Proud	980 l	2,908
5	BEERPONG - PIVO	219 l	681
6	IPA	134 l	332
7	Franz Spalek	126 l	285
8	KISELÁČ	24 l	55
9	ZON Lemonade	27 l	51
10	APA 12	1 l	3
	Total	19,793 l	53,990

Table 2.7: Most Consumed Beer Brands

The data demonstrate a strong preference for Radegast beer among festival attendees. The price may have influenced this, as the Radegast was generally slightly less expensive than, for instance, the Pilsner Urquell.

Key Takeaways

- The most consumed beer brand was **Radegast** with a total of **10,366** liters consumed.
- The second most consumed beer brand was the **Pilsner Urquell** with half the consumption of the Radegast.

Non-alcoholic Brands Analysis

This next section is bound to the non-alcoholic beverages, which were the second most consumed category.

Research Question

RQ17: What was the most popular non-alcoholic beverage brand, and how much was sold?

Aggregated similarly as in the previous analysis, the results in the Table 2.8 below show the distribution of the most consumed non-alcoholic brands during the event.

	Non-alcoholic brand	Total consumption	Units sold
1	Birell	2,186 l	4,893
2	ZON Lemonade	1,878 l	4,508
3	Water	984 l	2,855
4	Redbull	862 l	3,973
5	Lemonade	791 l	2,063
6	Coffee	204 l	1,855
7	Tea	54 l	153
8	Others	12 l	187
	Total	6,971 l	20,487

Table 2.8: Most Consumed Non-alcoholic Brands

The results indicate a close rivalry between the **Birell** and **ZON Lemonade** brands, with Birell emerging as the most consumed non-alcoholic beverage at the event. A total of **6,971** liters drank and **20,487** units sold established it as the most popular non-alcoholic beverage, commanding a **31.35%** share of the total consumption.

Key Takeaways

- The most consumed non-alcoholic brand was **Birell** with a total of **2,186** liters consumed.
- The second most consumed non-alcoholic brand was the **ZON Lemonade**.

Alcoholic Brands Analysis

The last part of this section focuses on other alcoholic beverages, such as spirits, shots, cocktails, and other alcoholic beverages.

Research Question

RQ16: What was the most popular brand of other alcoholic beverages, and how much was consumed and sold?

This analysis consisted of a total of 23 different brands consisting mostly of spirits and shots. Results presented in this section may be more biased than the previous ones, as the brand determination was more difficult due to the lack of consistent labeling. Moreover, previously it was necessary to classify the products with the volume information, which is not so straightforward for shots and spirits.

The results in the Table 2.9 indicate that the most consumed alcoholic brand was the **Beefeater** with a total of **966** liters consumed and **4,628** units sold. However, the second most consumed brand was **Absolut Vodka** with only **335** liters consumed, but with a much higher units sold count of **9,177**.

	Alcoholic brand	Total consumption	Units sold
1	Beefeater	966 l	4,628
2	Absolut Vodka	335 l	9,177
3	Jack Daniels	312 l	2,391
4	Becherovka Lemond	96 l	4,941
5	Havana	67 l	1,847
6	Malibu + Pineapple	55 l	707
7	Malfy Rosa + Tonic	45 l	217
8	Fire Štrúdl	28 l	770
9	Aperol	17 l	414
10	Others	59 l	1,872
	Total	1,980 l	26,964

Table 2.9: Most Consumed Alcoholic Brands

Key Takeaways

- The most consumed alcoholic brand was the **Beefeater** with a total of **966** liters consumed.
- The second most consumed alcoholic brand was the **Absolut Vodka** with a total of **335** liters consumed.

2.3.4 Summary

Insights into beverage consumption, covered in this section, should provide a better understanding of the overall preferences and total consumptions during the festival. As this data was not previously available, it was a valuable addition to the analysis and will play a significant role when presenting the results to the festival organizers.

The results showed the total consumption of all beverages, the insightful view on returnable cups usage, and the most consumed beverage brands in the most popular categories.

With the data available, there are still many more questions that could be answered, particularly in combination with customer data. However, clear bounds were set for this analysis to answer the most important questions asked by the festival organizers.

The next section will focus on the customer analysis to get a better understanding of the festival visitors and their behavior.

2.4 Customer Analysis

This section will focus on customer analysis, which should provide interesting insights into festival attendees' behavior and segmentation.

Because the organizers lacked deeper insights into festival attendees, this analysis should contribute significantly to a better understanding of the event.

The analysis addresses the remaining twelve research questions, which have been reorganized into four logical groups to improve the narrative flow:

1. Event Attendance and Timeline Analysis

- **RQ18:** *What was the total attendance at the event, and how many active customers were there each day?*
- **RQ21:** *What was the course of the event in terms of new visitors? And when were the largest “peaks”?*
- **RQ22:** *What is the average time of a visitor from arrival to first transaction?*
- **RQ23:** *What was the course of the event in terms of topping up credit on-site? And when were the largest “peaks”?*

2. Customer Segmentation

- **RQ27:** *How many customers added credit in advance online?*
- **RQ29:** *How many customers use the mobile app?*
- **RQ28:** *What was the distribution of customers by type (on-site, online, staff, guest, VIP)?*

3. Payment Behavior Analysis

- **RQ19:** *What is the distribution of bank cards used to refund credit?*
- **RQ20:** *W*
- **RQ24:** *How many customers topped up their credit multiple times, and how much only once?*

4. and Purchase Pattern Analysis

- **RQ25:** *What were the beverage preferences throughout the day?*
- **RQ26:** *What were the most common product combinations?*

2.4.1 Event Attendance and Timeline Analysis

This section should provide information about attendees’ initial behavior, total attendance throughout the festival, and subsequent top-up behavior.

Total Attendance and Daily Activity

Research Question

RQ18: What was the total attendance at the event, and how many active customers were there each day?

The analysis shows that the festival attracted a total of **10,009** unique customers over the three-day duration. However, the daily attendance figures show interesting patterns in how these customers were distributed throughout the festival days.

Looking at the daily active customers:

- Day 1 (Thursday): **6,214** active customers
- Day 2 (Friday): **5,832** active customers
- Day 3 (Saturday): **8,066** active customers

These numbers indicate that many customers attended multiple days of the festival, as the sum of daily attendees (**20,112**) is significantly higher than the unique customer count. This was expected, as the festival was designed to attract visitors for multiple days.

The attendance peaked on the final day of the festival, with Day 2 showing slightly lower attendance than the opening day. The significant increase in attendance on Day 3 suggests that the festival successfully attracted weekend visitors.

Key Takeaways

- Total unique customers over the three days: **10,009**.
- The highest attendance was on Day 3 (Saturday) with **8,066** active customers.
- The festival maintained consistent attendance on weekdays (Days 1–2).
- Significant increase in attendance (38% higher than Day 2) for the weekend (Day 3).

Visitor Arrival Patterns

Research Question

RQ21: What was the course of the event in terms of new visitors? And when were the largest “peaks”?

The analysis of visitor arrival patterns reveals distinct peak periods across the

three festival days, with significant variations in the rate of new visitor registrations throughout each day.




Figure 2.13: Visitor Arrival Patterns

Source: *Own creation*

The data reveals three distinct patterns across the festival days:

- **Day 1 (Thursday):** A rapid peak of approximately **1,873** new visitors was observed during the 16:00 hour, which was the time of the highest registration activity.

The day showed a clear pattern of increasing registrations from 14:00 to 16:00, followed by a gradual decline.

- **Day 2 (Friday):** A modest peak of **215** visitors at 15:00, but significantly lower new registrations than on Day 1.

This decrease in the number of new registrations was expected, as a significant number of visitors had already completed the registration process on Day 1.

- **Day 3 (Saturday):** The weekend attendees' influx was clear in the resurgence of new registrations, which reached a peak of **728** new visitors at 14:00.

The data indicates that the majority of visitor registrations occurred during the afternoon hours (14:00–17:00) on all days. Minimal new registrations were consistently observed during the early morning hours (before 10:00) and late evening hours (after 20:00).

Key Takeaways

- Highest single-hour registration peak: **1,873** visitors (Day 1, 16:00).
- The registration window that was most effective was from 14:00 to 17:00 on all days.
- Day 1 accounted for approximately **64.39%** of total registrations.
- Weekend (Day 3) saw renewed registration activity with peak of **728** new visitors.

Time to First Transaction

Research Question

RQ22: What is the average time of a visitor from arrival to first transaction?

The analysis reveals significant variations in how quickly different types of visitors made their first transaction after arrival. Overall, the average time to the first transaction was **80** minutes. However, this number alone is misleading, as evidenced by the much lower median of **7** minutes and mode of **3** minutes.




Figure 2.14: Time from Arrival to First Transaction Analysis

Source: *Own creation*

Breaking down the analysis by visitor type¹, reveals distinct patterns:

- **Regular:** Quickest to transact, with an average of **45** minutes and a median of only **6** minutes

¹Visitor or rather, chip types are described in the **Chip types enumeration** table.

- **Online:** Similar efficiency with an average of **67** minutes and a median of **7** minutes
- **Guests:** Took longer, averaging **141** minutes with a median of **10** minutes
- **Staff:** Showed significantly different behavior, with an average of **584** minutes and a median of **234** minutes

Visitor Type	Average	Mode	Median	Max
■ Guest	141 min	0 min	10 min	3,176 min
■ Regular	45 min	0 min	6 min	2,852 min
■ Online	67 min	3 min	7 min	3,199 min
■ Staff	584 min	7 min	234 min	3,492 min
Overall	80 min	3 min	7 min	3,492 min

Table 2.10: Time from Arrival to First Transaction

Source: *Own creation*

The significant difference between average and median times across all categories suggests a right-skewed distribution, implying that while most visitors completed their first transaction quickly, others took much longer. This was especially true for regular attendees, who would sometimes arrive early to check in and then return later to make their first purchase.

Staff members who were not there primarily to consume, on the other hand, exhibited a different behavior pattern, with their first transaction occurring later. This was most likely due to their responsibilities at the event, which may have prevented them from making purchases during working hours.

Key Takeaways

- Most visitors (as indicated by the median) made their first transaction within **7** minutes of arrival.
- Regular visitors were the most efficient, typically transacting within **6** minutes.
- Staff members showed distinctly different behavior, likely due to their different role at the event.
- The large gap between mean and median times suggests some visitors waited significantly longer than others before their first transaction.

Credit Top-up Patterns

Research Question

RQ23: What was the course of the event in terms of topping up credit on-site?
And when were the largest “peaks”?

The examination of credit top-up patterns reveals distinct peak periods and daily patterns throughout the event. A total of **17,233** top-up transactions were processed, with activity varying significantly throughout the day.




Figure 2.15: Credit Top-up Patterns During Event

Source: *Own creation*

The data shows three major peak periods:

- **Day 1 Peak:** The highest single-hour volume occurred at 16:00 with **1,793** top-ups
- **Day 2 Peak:** A moderate peak at 17:00 with **672** top-ups
- **Day 3 Peak:** A strong resurgence with **1,162** top-ups at 16:00

Each day had a similar pattern, with activity increasing in the early afternoon, peaking in the late afternoon (between 16:00 and 18:00), and gradually decreasing into the evening. The overnight hours (00:00–12:00) saw minimal top-up activity.

Key Takeaways

- Peak top-up activity consistently occurred during late afternoon hours
- Day 1 saw the highest single-hour volume with **1,793** top-ups
- Clear daily pattern of afternoon peaks and overnight lulls
- Day 3 showed sustained high activity with multiple hours exceeding **800** top-ups

2.4.2 Customer Segmentation

TODO: Distribution and characteristics of customer types, Analysis of online vs on-site customer behavior, Digital service adoption (mobile app, online top-ups)

Research Question

RQ27: How many customers added credit in advance online?

Research Question

RQ29: How many customers use the mobile app?

Research Question

RQ28: What was the distribution of customers by type (on-site, online, staff, guest, VIP)?

2.4.3 Payment Behavior Analysis

TODO: Analysis of payment methods and card schemes, Credit top-up frequency patterns, Duration-based analysis of customer spending

Research Question

RQ19: What is the distribution of bank cards used to refund credit?

Research Question

RQ20: W

Research Question

RQ24: How many customers topped up their credit multiple times, and how much only once?

2.4.4 Purchase Pattern Analysis

TODO: Temporal analysis of drink preferences, Product combination analysis, Purchase behavior insights

Research Question

RQ25: What were the beverage preferences throughout the day?

Research Question

RQ26: What were the most common product combinations?

3 Dashboard Implementation

TODO: This chapter describes the implementation of the analytical dashboard using Dash and Plotly.

TODO: Tools and technologies used - Dash, Plotly, Python, PostgreSQL, SQLAlchemy, Pandas, Mantine Dash Components etc.

TODO: Describe async loops for data fetching and Dash background callbacks, f*cking issues with diskcache DiskcacheManager issues with SQLite concurrent connections, describe how we tried for 12hours straight to fix it and switch to celery + redis for async tasks but then switched back to diskcache with one-liner fix, fml

TODO: Describe DB query management - QueryDefinition, QueryRegistry, QueryManger, QueryParameter and other internal classes

TODO: Describe additional simple FS query caching mechanism

TODO: Describe process of implementing individual dashboard sections, structure and most importantly the chosen visualizations/presentation of data

TODO: Defend why we didn't go for deployment to cloud – sorry not sorry I'm not a lunatic, hate Python env and deployment issues

TODO: Final dashboard presentation and features

TODO: Next steps and missing features and future improvements we would do if we had more time

4 Conclusion

4.1 Summary of Work

TODO: Write a summary of the work done in this thesis, recap the main objectives and how they were achieved. Note the key contributions of the research and practical side.

4.2 Reflections

TODO: Write reflections on the project, what was learned, how the project could change festival operations, etc.

4.3 Future Directions

TODO: Write about the future directions of the project, what could be done next, what could be improved, what could be added, etc.

List of Figures

1.1	Knowledge Data Discovery workflow diagram.	33
2.1	Top-Up Transactions by Payment Method	36
2.2	Sales of the Organizer vs. External Vendors	37
2.3	Remaining Chip Balances Sankey Diagram	39
2.4	Breakdown of Direct Revenue Streams	41
2.5	Breakdown of All Revenue Streams	42
2.6	Overall Cash Flow Diagram	43
2.7	Transactions Processing Peaks	46
2.8	Transaction Processing Peaks	47
2.9	Best Top-Up Points	49
2.10	Best Products by Category	54
2.11	Returnable Cups	56
2.12	Most Consumed Beer Brands	59
2.13	Visitor Arrival Patterns	65
2.14	Time from Arrival to First Transaction Analysis	67
2.15	Credit Top-up Patterns During Event	69

Note: Unless stated otherwise, all figures are the author's own work.

List of Tables

1.1	Product categories.	28
1.2	Customer chips function return table.	30
1.3	Chip types enumeration	31
2.1	Revenue Summary Breakdown	42
2.2	Best Top-Up Points	49
2.3	Best Sale Points	51
2.4	Best Vendors	52
2.5	Best Products	53
2.6	Total Beverage Consumption	57
2.7	Most Consumed Beer Brands	59
2.8	Most Consumed Non-alcoholic Brands	60
2.9	Most Consumed Alcoholic Brands	61
2.10	Time from Arrival to First Transaction	67

List of Source Codes

1	Anonymization configuration example.	25
---	--	----

List of Acronyms

List of Appendices

Appendix A Source code of the application